

Welcome to CSE 502

Introduction & Review

Today's Lecture

- Course Overview
- Course Topics
- Grading
- Logistics
- Academic Integrity Policy
- Homework #1 (of 1)
- Quiz #1 (of 1)
- Key concepts from Undergrad Computer Architecture

Course Overview (1/3)

- Caveat 1: I'm new here.
- Caveat 2: This is a new course.
- Computer Architecture is
 - ... the ***science*** and ***art*** of selecting and interconnecting hardware and software components to ***create computers*** ...

Course Overview (2/3)

- Ever wonder what's inside that box, anyway?
- Computer Architecture is an umbrella term
 - Architecture: software-visible interface
 - Micro-architecture: internal organization of components
- This course is mostly about ***micro-architecture***
 - What's inside the processor (CPU)
 - What implications this has on software

Course Overview (3/3)

- This course is hard, roughly like CSE 506
 - In CSE 506, you learn what's inside an OS
 - In CSE 502, you learn what's inside the CPU
- This is a project course
 - Learn why things are the way they are, first hand
 - We will “build” emulators of CPU components
 - If you don't know C++, you need to learn it quickly

Course Topics

- Intro/Review
- ***Memory Hierarchy***
- Pipelining
- Instruction-Level Parallelism
- ***Processor Front-end***
- ***Execution Core***
- Multi-[socket(SMP,DSM) | thread(SMT,CMT) | core(CMP)]
- Prefetching
- Vector Processing and GPUs

Will devote most time to items in ***bold***

Grading

kk	Due Date	Points	Grading	Required?
1 Quiz	Today	0	Curve 0 to 100	Yes
1 Homework	Last class	10	Curve 0 to 100	No
1 Warm-up Project	Feb 25	20	See below	No
1 Course Project	Last class	100	See below	Yes
1 Final		30	Absolute value	No
Participation		10	Curve 0 to 100	No

Warm-up Project	Points	Course Project	Points
1 Port, Direct-Mapped, 16K	10	Single-cycle	50
2 Port, Direct-Mapped, 16K	12	5+ stage pipeline	60
1 Port, Set-associative, 16K	14	Super-scalar	70
2 Port, DM, 64K, Non-Block	16	Out-of-order	80
2 Port, SA, 64K, Non-Block	18	Out-of-order, Branch-Pred	90
2 Port, SA, 64K, NB, Way-pred.	20	Out-of-order, BP, SMT	100

Without curve, need 100 points to get an A

Logistics (1/3)

- Project milestones
 - There are **no** official project milestones
 - If **you** need milestones, send me a milestone schedule
 - I will deduct 5 points for each milestone you miss
- Book
 - Hennessy & Patterson, use it as a reference
 - Exam questions will mostly come from the book
 - Appendices as important as the book chapters

Logistics (2/3)

- Working in groups
 - Permitted on everything except Quiz and Final
 - Groups may range in size from 1 to 72 people
- Attendance
 - Optional (but highly advised)
 - No laptop, tablet, or phone use in class
 - Don't test me - I **will deduct** grade points

Logistics (3/3)

- Blackboard
 - This will be my first time using it, don't expect much there
- Course Mailing List
 - Subscription is **required**
<http://lists.cs.stonybrook.edu/mailman/listinfo/cse502>
- Quiz
 - Completion is **required**
 - If you missed the 1st class, come to office hours for it

Academic Integrity Policy

- Probably different from other classes
 - Much more open, but much more strict
 - Resembles the “real world”
- Actual Policy:
 - All submitted work must have an explicit Copyright label containing your name.
 - All submitted work must have an explicit license.
 - All Copyright laws of the United States must be respected.

Copyright Example 1

- Copyright © 2013 by Mike Ferdman.
All rights reserved.

Copying work labeled this way is prohibited

Copyright Example 2

- Copyright © 2013 by Mike Ferdman.
Permission to copy and distribute verbatim copies
permitted.

Copying is OK, but without ANY modifications

Copyright Example 3

- Copyright © 2013 by Mike Ferdman.
This work is licensed under GPLv3, details in accompanying COPYING file.

Copying is OK, but you must allow others to copy too

Questions?

Homework #1 (of 1) part 1

- Start with

```
volatile int x = 0;
int main(int argc, char* argv[]) {
    for(int y=0; y<1234567890; ++y) ++x;
    return x;
}
```

- Implement the ***most accurate*** way to measure time
 - How long does the `for()` loop take to run?

Quiz #1 (of 1)

Review

- Understanding and Measuring Performance
- Memory Locality
- Power and Energy
- Parallelism and Critical Paths
- Instruction Set Architecture
- Basic Processor Organization

This is intended to be a review, ***not*** an introduction!


Amdahl's Law

$$\text{Speedup} = \text{time}_{\text{without enhancement}} / \text{time}_{\text{with enhancement}}$$

An enhancement speeds up fraction f of a task by factor S


$$\text{time}_{\text{new}} = \text{time}_{\text{orig}} \cdot ((1-f) + f/S)$$

$$S_{\text{overall}} = 1 / ((1-f) + f/S)$$



The Iron Law of Processor Performance

$$\frac{\textit{Time}}{\textit{Program}} = \frac{\textit{Instructions}}{\textit{Program}} \times \frac{\textit{Cycles}}{\textit{Instruction}} \times \frac{\textit{Time}}{\textit{Cycle}}$$



Architects target CPI, but ***must*** understand the others

Performance

- Latency (execution time): time to finish one task
- Throughput (bandwidth): number of tasks/unit time
 - Throughput can exploit parallelism, latency can't
 - Sometimes complimentary, often contradictory
- Example: move people from A to B, 10 miles
 - Car: capacity = 5, speed = 60 miles/hour
 - Bus: capacity = 60, speed = 20 miles/hour
 - Latency: car = 10 min, bus = 30 min
 - Throughput: car = 15 PPH (count return trip), bus = 60 PPH

No right answer: pick metric for *your* goals

Performance Improvement

- Processor A is X times faster than processor B if
 - $\text{Latency}(P,A) = \text{Latency}(P,B) / X$
 - $\text{Throughput}(P,A) = \text{Throughput}(P,B) * X$
- Processor A is X% faster than processor B if
 - $\text{Latency}(P,A) = \text{Latency}(P,B) / (1+X/100)$
 - $\text{Throughput}(P,A) = \text{Throughput}(P,B) * (1+X/100)$
- Car/bus example
 - Latency? Car is 3 times (200%) faster than bus
 - Throughput? Bus is 4 times (300%) faster than car

Partial Performance Metrics Pitfalls

- Which processor would you buy?
 - Processor A: CPI = 2, clock = 2.8 GHz
 - Processor B: CPI = 1, clock = 1.8 GHz
 - Probably A, but B is faster (assuming same ISA/compiler)
- Classic example
 - 800 MHz Pentium III faster than 1 GHz Pentium 4
 - Same ISA and compiler

Averaging Performance Numbers (1/2)

- Latencies are additive, throughput is not
 $\text{Latency}(P1+P2,A) = \text{Latency}(P1,A) + \text{Latency}(P2,A)$
 $\text{Throughput}(P1+P2,A) \neq \text{Throughput}(P1,A) + \text{Throughput}(P2,A)$
- Example:
 - 1 mile @ 30 miles/hour + 1 mile @ 90 miles/hour
 - Average is **not 60** miles/hour
 - 0.033 hours at 30 miles/hour + 0.01 hours at 90 miles/hour
 - Average is **only 47** miles/hour! (2 miles / (0.033 + 0.01 hours))

Averaging Performance Numbers (2/2)

- Arithmetic: times
 - proportional to time
 - e.g., latency
- Harmonic: rates
 - inversely proportional to time
 - e.g., throughput
- Geometric: ratios
 - unit-less quantities
 - e.g., speedups

$$\frac{1}{n} \sum_{i=1}^n Time_i$$

$$\frac{n}{\sum_{i=1}^n \frac{1}{Rate_i}}$$

$$\sqrt[n]{\prod_{i=1}^n Ratio_i}$$

Memorize these to avoid looking them up later

Locality Principle

- Recent past is a good indication of near future



Temporal Locality: If you looked something up, it is very likely that you will look it up again soon

Spatial Locality: If you looked something up, it is very likely you will look up something nearby soon

Power vs. Energy (1/2)

- Power: instantaneous rate of energy transfer
 - Expressed in Watts
 - In Architecture, implies conversion of electricity to heat
 - $\text{Power}(\text{Comp1} + \text{Comp2}) = \text{Power}(\text{Comp1}) + \text{Power}(\text{Comp2})$
- Energy: measure of using power for some time
 - Expressed in Joules
 - $\text{power} * \text{time} \text{ (joules = watts * seconds)}$
 - $\text{Energy}(\text{OP1} + \text{OP2}) = \text{Energy}(\text{OP1}) + \text{Energy}(\text{OP2})$


Power vs. Energy (2/2)



Does this example help or hurt?

Why is energy important?

- Because electricity consumption has costs
 - Impacts battery life for mobile
 - Impacts electricity costs for tethered
 - Delivering power for buildings, countries
 - Gets worse with larger data centers (\$7M for 1000 racks)



Why is power important?

- Because Energy is $\text{Power} \times \text{Time}$
- Because of delivery, packaging, and cooling costs
 - Increased cost of thermal packing
 - \$1/W for CPUs > 35W
 - Temperature, noise issues ...
- Because of compaction, density, and reliability
 - Thermal failures
 - 50% server reliability degradation for +10oC
 - 50% decrease in hard disk lifetime for +15oC

Power: The Basics

- Dynamic power vs. Static power
 - Static: “leakage” power
 - Dynamic: “switching” power
 - Dynamic power dominates, but static increasing in importance
- Static power: steady, per-cycle energy cost
- Dynamic power: capacitive and short-circuit
 - Capacitive power: charging/discharging at $0 \rightarrow 1$ and $1 \rightarrow 0$
 - Short-circuit power: brief short-circuit during transitions

Dynamic Power Dissipation (Capacitive)

Capacitance:

Function of wire length,
transistor size

Supply Voltage:

Function of technology and
operating frequency


$$\text{Power} \approx \frac{1}{2} CV^2Af$$

Activity factor:

Average fraction of all possible
transitions ($0 \rightarrow 1$ and $1 \rightarrow 0$) per cycle?



Clock frequency:

Function of desired
performance


Lowering Dynamic Power

- Reducing Voltage (V) has quadratic effect
 - Has a negative (\sim linear) effect on frequency
 - Limited by technology (insufficient difference of 1 & 0)
- Lowering Capacitance (C) has linear effect
 - May improve frequency
 - Limited by technology (small transistors, short wires)
- Reducing switching Activity (A) has linear effect
 - A function of signal transition stats
 - Turns off idle units to reduce activity
 - Impacted by logic and architecture decisions


Leakage Power (1/3)



Leakage Power (2/3)



Leakage Power (3/3)



(important take-away is on the next slide)

Thermal Runaway

- Leakage is a function of temperature
- ↑ Temp leads to ↑ Leakage
- Which burns more power
- Which leads to ↑ Temp, which leads to...


$$I_{\text{sub}} = K_1 W e^{-V_{\text{th}}/nV_{\theta}} (1 - e^{-V/V_{\theta}})$$

Positive feedback loop will melt your chip

Power Management in Processors


- Clock gating

- Stop switching in unused components
- Done automatically in most designs
- Near instantaneous on/off behavior



- Power gating

- Turn off power to unused cores/caches
- High latency for on/off
 - Saving SW state, flushing dirty cache lines, turning off clock tree
 - Carefully done to avoid voltage spikes or memory bottlenecks
- Issue: Area & power consumption of power gate
- Opportunity: use thermal headroom for other cores



DVFS: Dynamic Voltage/Frequency Scaling

- Set frequency to the lowest needed
 - Execution time = $IC * CPI * F$
- Scale back V to lowest for that frequency
 - Lower voltage \rightarrow slower transistors
 - Power $\approx C * V^2 * F$
- Provides P states for power management
 - Heavy load: frequency, voltage, power high
 - Light load: frequency, voltage, power low
 - Trade-off: power savings vs overhead of scaling
 - Effectiveness limited by voltage range

Parallelism: Work and Critical Path

- Parallelism: number of independent tasks available
- Work (T_1): time on sequential system
- Critical Path (T_∞): time on infinitely-parallel system

- Average Parallelism:


$$P_{\text{avg}} = T_1 / T_\infty$$

- For a p -wide system:

$$T_p \geq \max\{T_1/p, T_\infty\}$$

$$P_{\text{avg}} \gg p \Rightarrow T_p \approx T_1/p$$

```
x = a + b;
y = b * 2
z = (x - y) * (x + y)
```



Can trade off frequency for parallelism

ISA: A contract between HW and SW

- ISA: Instruction Set Architecture
 - A well-defined hardware/software interface
- The “contract” between software and hardware
 - Functional definition of operations supported by hardware
 - Precise description of how to invoke all features
- No guarantees regarding
 - How operations are implemented
 - Which operations are fast and which are slow (and when)
 - Which operations take more energy (and which take less)

Components of an ISA

- Programmer-visible states
 - Program counter, general purpose registers, memory, control registers
- Programmer-visible behaviors
 - What to do, when to do it

**Example “register-transfer-level”
description of an instruction**

```
if imem[pc]==“add rd, rs, rt”  
then
```

```
    pc  $\leftarrow$  pc+1
```

```
    gpr[rd]=gpr[rs]+gpr[rt]
```

- A binary encoding

ISAs last forever, don't add stuff you don't need

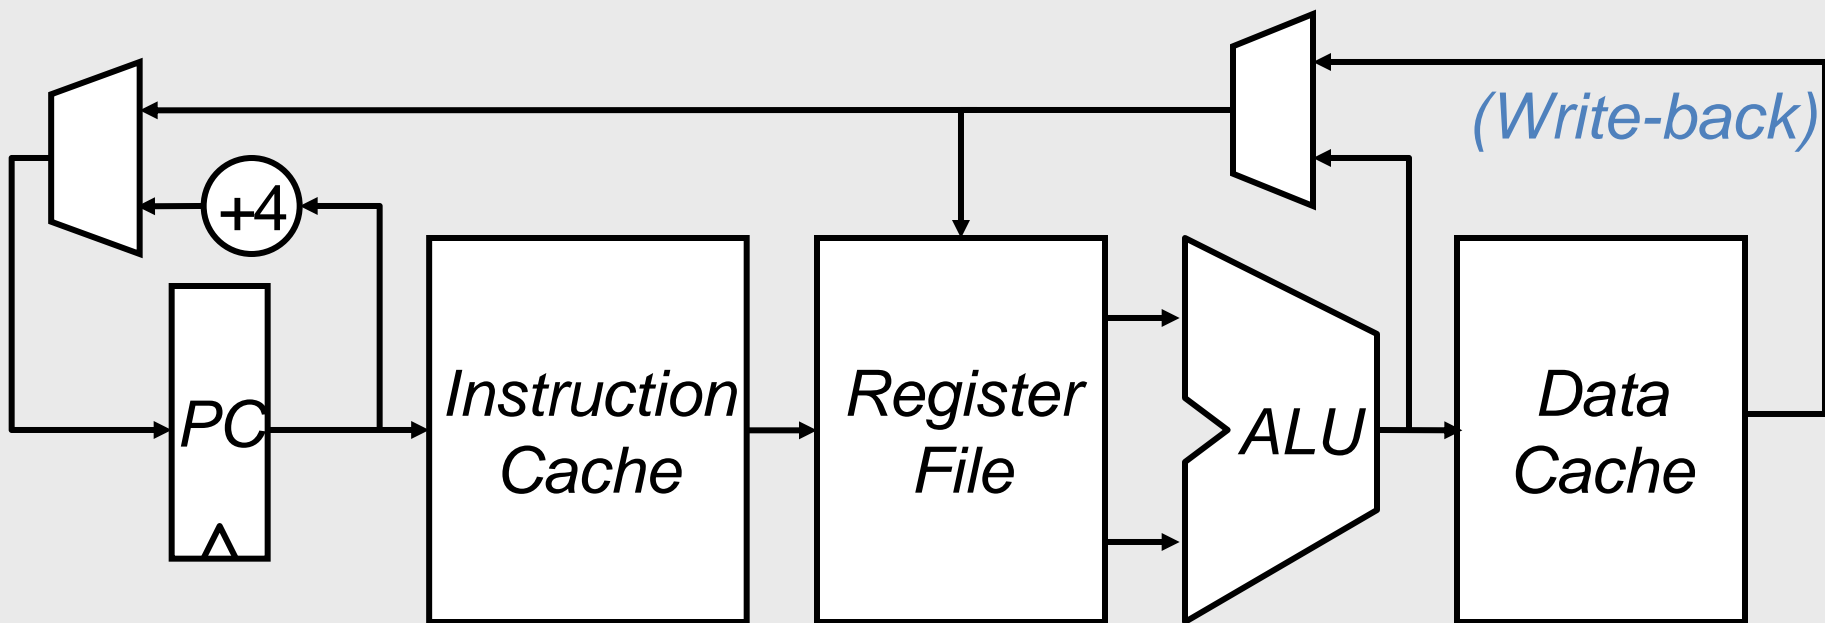
RISC vs. CISC

- Recall Iron Law:
 - $(\text{instructions/program}) * (\text{cycles/instruction}) * (\text{seconds/cycle})$
- CISC (Complex Instruction Set Computing)
 - Improve “instructions/program” with “complex” instructions
 - Easy for assembly-level programmers, good code density
- RISC (Reduced Instruction Set Computing)
 - Improve “cycles/instruction” with many single-cycle instructions
 - Increases “instruction/program”, but hopefully not as much
 - Help from smart compiler
 - Perhaps improve clock cycle time (seconds/cycle)
 - via aggressive implementation allowed by simpler instructions

Today's x86 chips translate CISC into ~RISC

Prototypical Processor Organization

Addr-gen. Fetch Decode Issue Execute Memory



Conclusion

- Know the topics from the Review section
 - If you don't, you need to catch up
- Make sure you know C++
 - If you don't, get a book today
- So far, we had intro + review potpourri
- The rest of this course will be very ***unlike*** this lecture
 - Very few new terms
 - Practically no formulas
 - Lots of new material

Questions?