

CSE 502:

Computer Architecture

Instruction Decode

RISC ISA Format

- This should be review...
 - Fixed-length
 - MIPS all insts are 32-bits/4 bytes
 - Few formats
 - MIPS has 3 formats: R (reg, reg, reg), I (reg, reg, imm), J (addr)
 - Alpha has 5: Operate, Op w/ Imm, Mem, Branch, FP
 - Regularity across formats (when possible/practical)
 - MIPS & Alpha opcode in same bit-position for all formats
 - MIPS rs & rt fields in same bit-position for R and I formats
 - Alpha ra/fa field in same bit-position for all 5 formats

RISC Decode (MIPS)



000xxx = Br/Jump
(except for 000000)

opcode[2,0]

			010	011	100	101	110	111
000	<i>func</i>	<i>rt</i>	j	jal	beq	bne	blez	bgtz
001	addi	addiu	slti	sltiu	andi	ori	xori	lui
010				rs				
011								
100	lb	lh	lwl	lw	lbu	lhu	lwr	
101	sb	sh	swl	sw			swr	
110	lwc0	lwc1	lwc2	lwc3				
111	swc0	swc1	swc2	swc3				

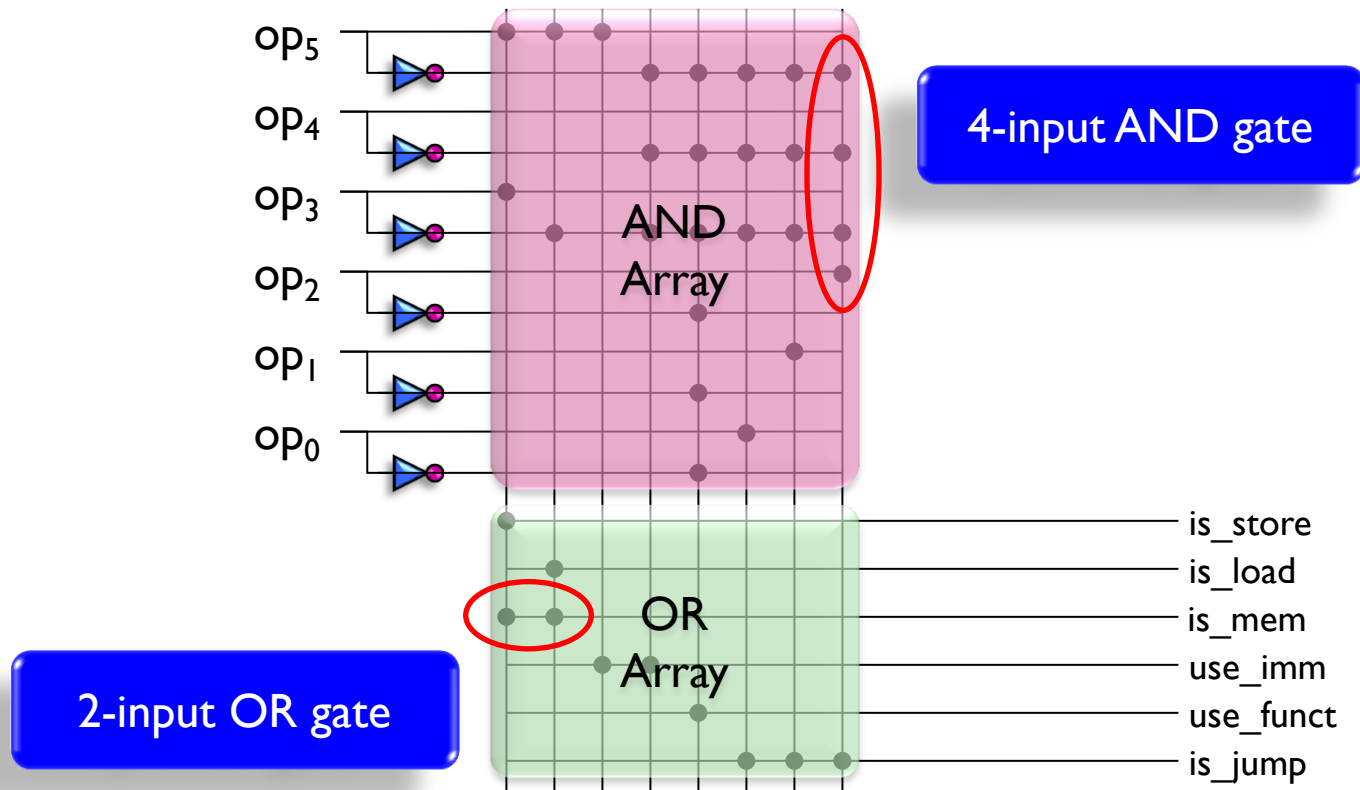
001xxx = Immediate

xxxxx = Memory
(1x0:LD, 1x1:ST)

PLA Decoders (1/2)

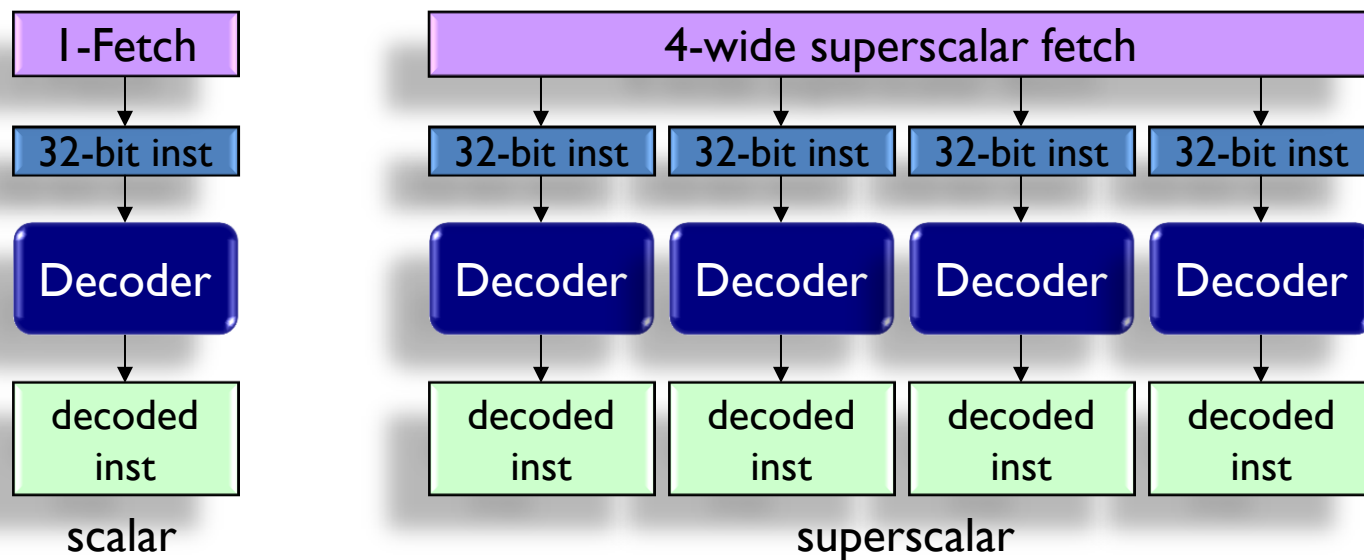
- PLA = Programmable Logic Array
- Simple logic to transform opcode to control signals
 - $\text{is_jump} = \neg \text{op5} \ \& \ \neg \text{op4} \ \& \ \neg \text{op3} \ \& \ (\text{op2} \mid \text{op1} \mid \text{op0})$
 - $\text{use_funct} = \neg \text{op5} \ \& \ \neg \text{op4} \ \& \ \neg \text{op3} \ \& \ \neg \text{op2} \ \& \ \neg \text{op1} \ \& \ \neg \text{op0}$
 - $\text{use_imm} = \text{op5} \mid \neg \text{op5} \ \& \ \neg \text{op4} \ \& \ \text{op3}$
 - $\text{is_load} = \text{op5} \ \& \ \neg \text{op3}$
 - $\text{is_store} = \text{op5} \ \& \ \text{op3}$

PLA Decoders (2/2)



Superscalar Decode for RISC ISAs

- Decode X insns. per cycle (e.g., 4-wide)
 - Just duplicate the hardware
 - Instructions aligned at 32-bit boundaries



CISC ISA

- RISC focus on fast access to information
 - Easy decode, I\$, large RF's, D\$
- CISC focus on max expressiveness per min space
 - Designed in era with fewer transistors, chips
 - Each memory access very expensive
 - Pack as much work into as few bytes as possible
 - More “expressive” instructions
 - Better potential code generation in theory
 - More complex code generation in practice

ADD in RISC ISA

Mode	Example	Meaning
Register	ADD R4, R3, R2	$R4 = R3 + R2$

ADD in CISC ISA

Mode	Example	Meaning
Register	ADD R4, R3	$R4 = R4 + R3$
Immediate	ADD R4, #3	$R4 = R4 + 3$
Displacement	ADD R4, 100(R1)	$R4 = R4 + \text{Mem}[100+R1]$
Register Indirect	ADD R4, (R1)	$R4 = R4 + \text{Mem}[R1]$
Indexed/Base	ADD R3, (R1+R2)	$R3 = R3 + \text{Mem}[R1+R2]$
Direct/Absolute	ADD R1, (1234)	$R1 = R1 + \text{Mem}[1234]$
Memory Indirect	ADD R1, @(R3)	$R1 = R1 + \text{Mem}[\text{Mem}[R3]]$
Auto-Increment	ADD R1, (R2)+	$R1 = R1 + \text{Mem}[R2]; R2++$
Auto-Decrement	ADD R1, -(R2)	$R2--; R1 = R1 + \text{Mem}[R2]$


x86

- CISC, stemming from the original 4004 (~1971)
- Example: “Move” instructions
 - General Purpose data movement
 - $R \rightarrow R$, $M \rightarrow R$, $R \rightarrow M$, $I \rightarrow R$, $I \rightarrow M$
 - Exchanges
 - $EAX \leftrightarrow ECX$, byte order within a register
 - Stack Manipulation
 - $\text{push pop } R \leftrightarrow \text{Stack}$, PUSH/POPA
 - Type Conversion
 - Conditional Moves

Many ways to do the same/similar operation

x86 Encoding


- Basic x86 Instruction:





- Opcode has flag indicating Mod R/M is present
 - Most instructions use the Mod R/M byte
 - Mod R/M specifies if optional SIB byte is used
 - Mod R/M and SIB may specify additional constants

Instruction length not known until *after* decode

x86 Mod R/M Byte



- Mode = 00: No-displacement, use Mem[Rmmm]
- Mode = 01: 8-bit displacement, Mem[Rmmm+disp])]
- Mode = 10: 32-bit displacement (similar to previous)
- Mode = 11: Register-to-Register, use Rmmm

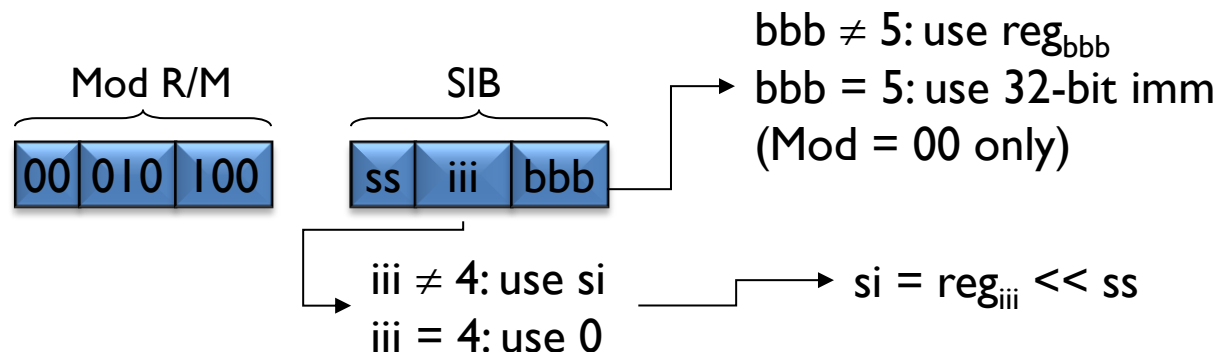


x86 Mod R/M Exceptions

- Mod=00, R/M = 5 \rightarrow get operand from 32-bit imm
 - Add


00	010	101	0cff1234
----	-----	-----	----------

 $EDX = EDX + \text{Mem}[0cff1234]$
- Mod=00, 01 or 10, R/M = 4 \rightarrow use the “SIB” byte
 - SIB = Scale/Index/Base




x86 Opcode Confusion

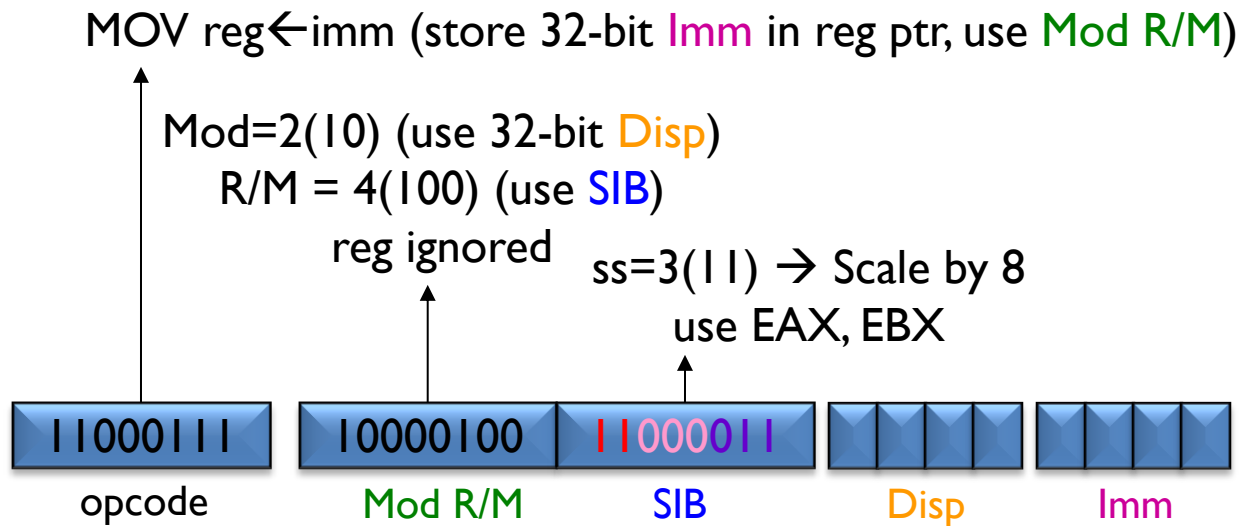
- There are different opcodes for $A \leftarrow B$ and $B \rightarrow A$



- If Opcode = 0F, then use next byte as opcode
- If Opcode = D8-DF, then FP instruction



x86 Decode Example



$$* ((EAX \ll 3) + EBX + Disp) = Imm$$

Total: 11 byte instruction

Note: Add 4 prefixes, and you reach the max size

RISC (MIPS) vs CISC (x86)

lui R1, Disp[31:16]

ori R1, R1, Disp[15:0]

add R1, R1, R2

shli R3, R3, 3

add R3, R3, R1

lui R1, Imm[31:16]

ori R1, R1, Imm[15:0]

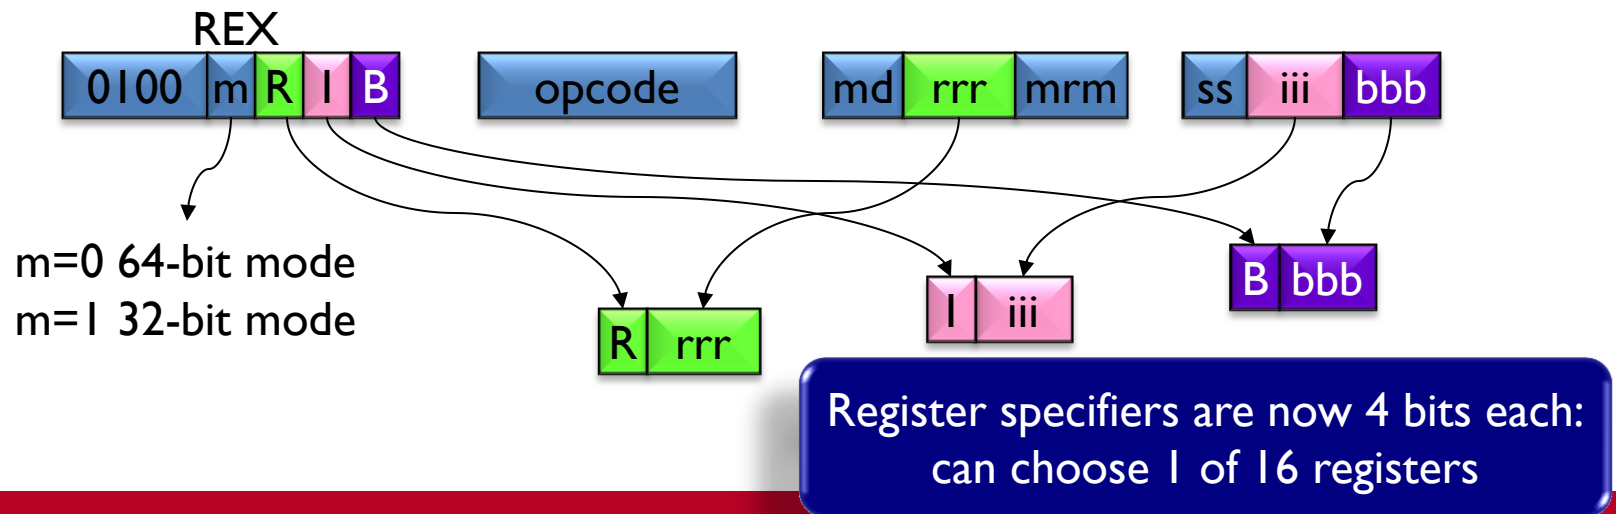
st [R3], R1


MOV [EBX+EAX*8+Disp], Imm

8 insns. at 32 bits each **vs** 1 insn. at 88 bits: **2.9x!**

x86-64 / EM64T

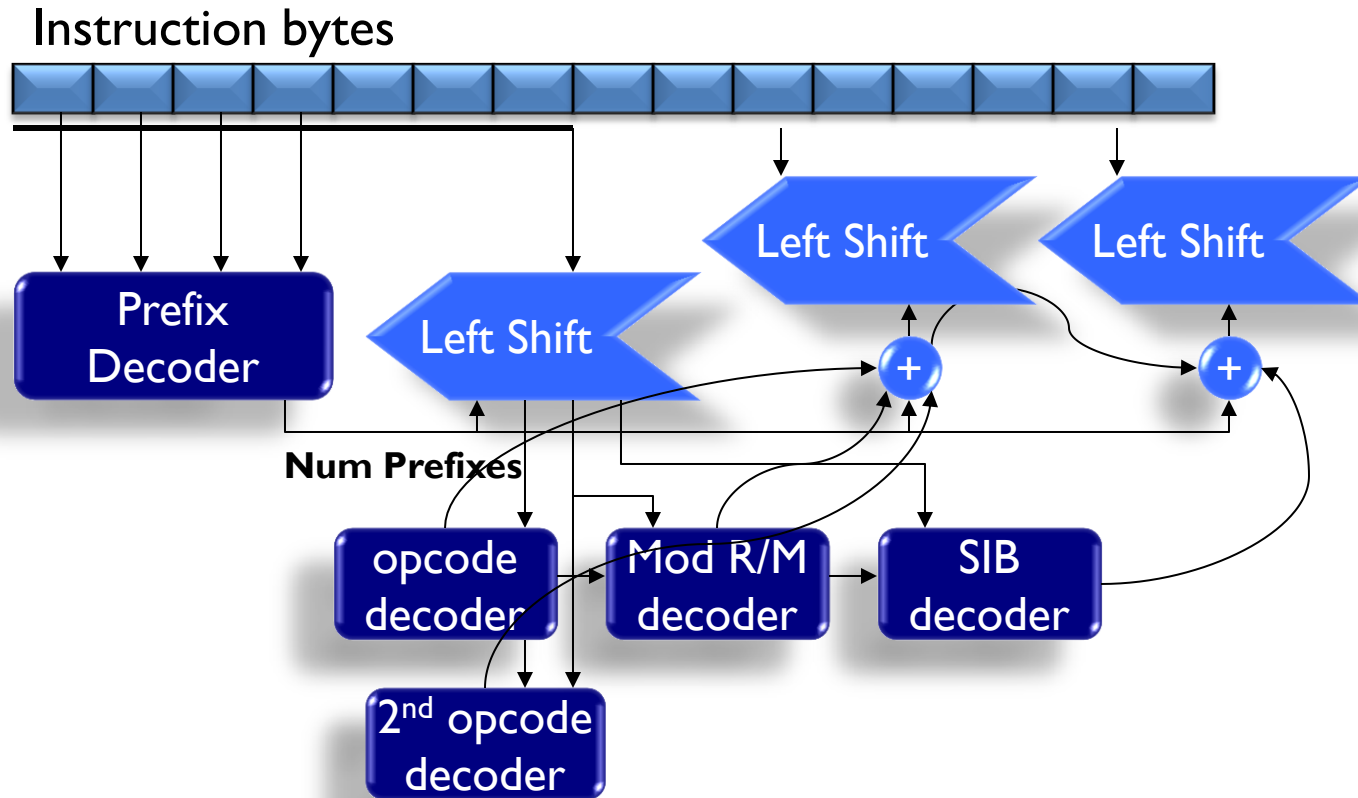
- 8→16 general purpose registers
 - But we only used to have 3-bit register fields...
- Registers extended from 32→64 bits each
- Default: instructions still 32-bit
 - New “REX” prefix byte to specify additional information





Ugly? Scary? But it works...

x86 Decode Hardware



Decoded x86 Format

- RISC: easy to expand → union of needed info
 - Generalized opcode (not too hard)
 - Reg1, reg2, reg3, immediate (possibly extended)
 - Some fields ignored
- CISC: union of all possible info is huge
 - Generalized opcode (too many options)
 - Up to 3 regs, 2 immediates
 - Segment information
 - “rep” specifiers
 - Would lead to 100’s of bits
 - Common case only needs a fraction → a lot of waste

Too expensive to decode x86 into control bits

x86 → RISC-like mops

- x86 decoded into “uops”(Intel) or “ROPs”(AMD)
... (micro-ops or RISC-ops)
 - Each uop is RISC-like
 - uops have limitations to keep union of info practical

ADD EAX, EBX	————→	ADD EAX, EBX	1 uop
ADD EAX, [EBX]	————→	Load tmp = [EBX] ADD EAX, tmp	2 uops
ADD [EAX], EBX	————→	Load tmp = [EAX] ADD tmp, EBX STA EAX STD tmp	4 uops

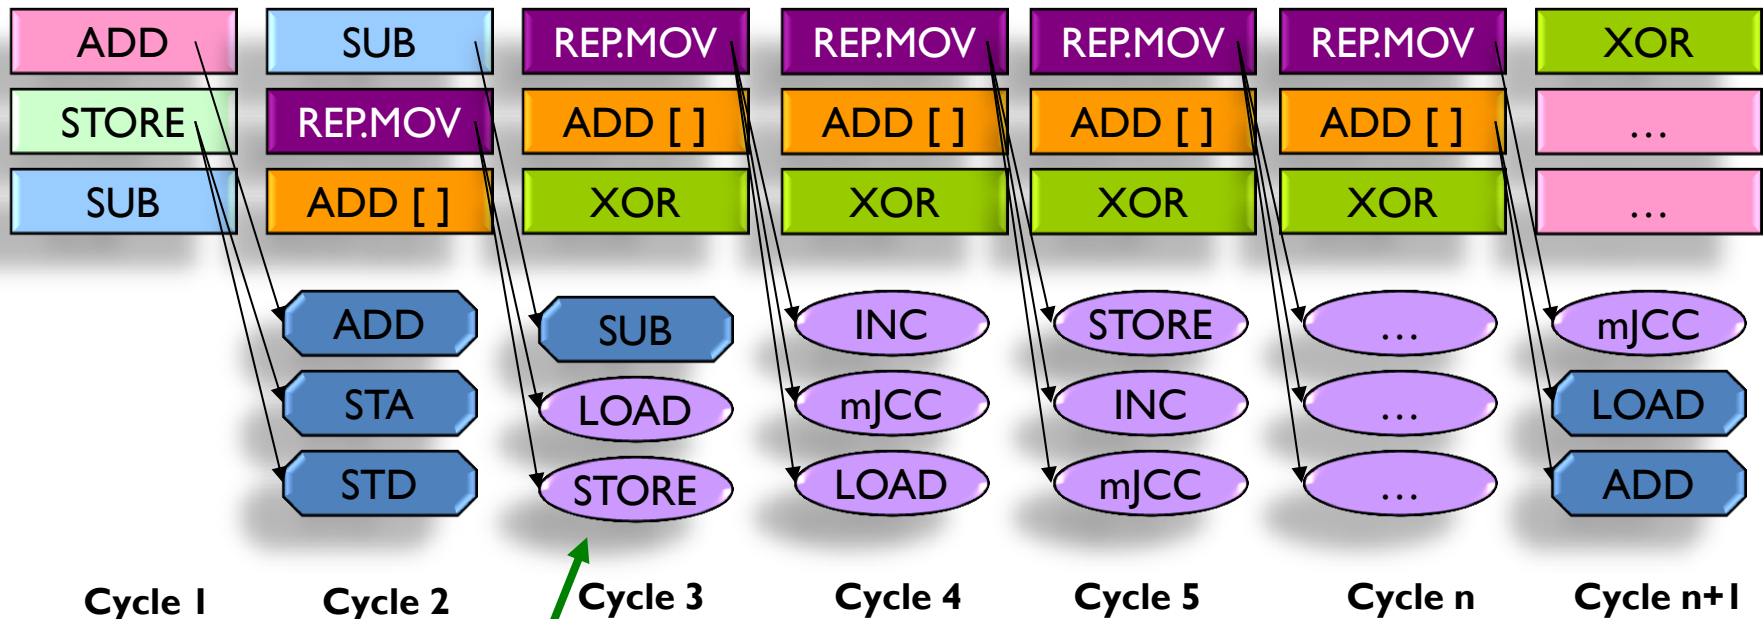
uop Limits

- How many uops can a decoder generate?
 - For complex x86 insts, many are needed (10's, 100's?)
 - Makes decoder horribly complex
 - Typically there's a limit to keep complexity under control
 - One x86 instruction → 1-4 uops
 - Most instructions translate to 1.5-2.0 uops
- What if a complex insn. needs more than 4 uops?


UROM/MS for Complex x86 Insts

- UROM (microcode-ROM) stores the uop equivalents
 - Used for nasty x86 instructions
 - Complex like > 4 uops (PUSHA or STRREP.MOV)
 - Obsolete (like AAA)
- Microsequencer (MS) handles the UROM interaction

UROM/MS Example (3 uop-wide)



Complex instructions, get uops from mcode sequencer



Superscalar CISC Decode

- Instruction Length Decode (ILD)
 - Where are the instructions?
 - Limited decode – just enough to parse prefixes, modes
- Shift/Alignment
 - Get the right bytes to the decoders
- Decode
 - Crack into uops

Do this for N instructions per cycle!

ILD Recurrence/Loop

$$PC_i = X$$

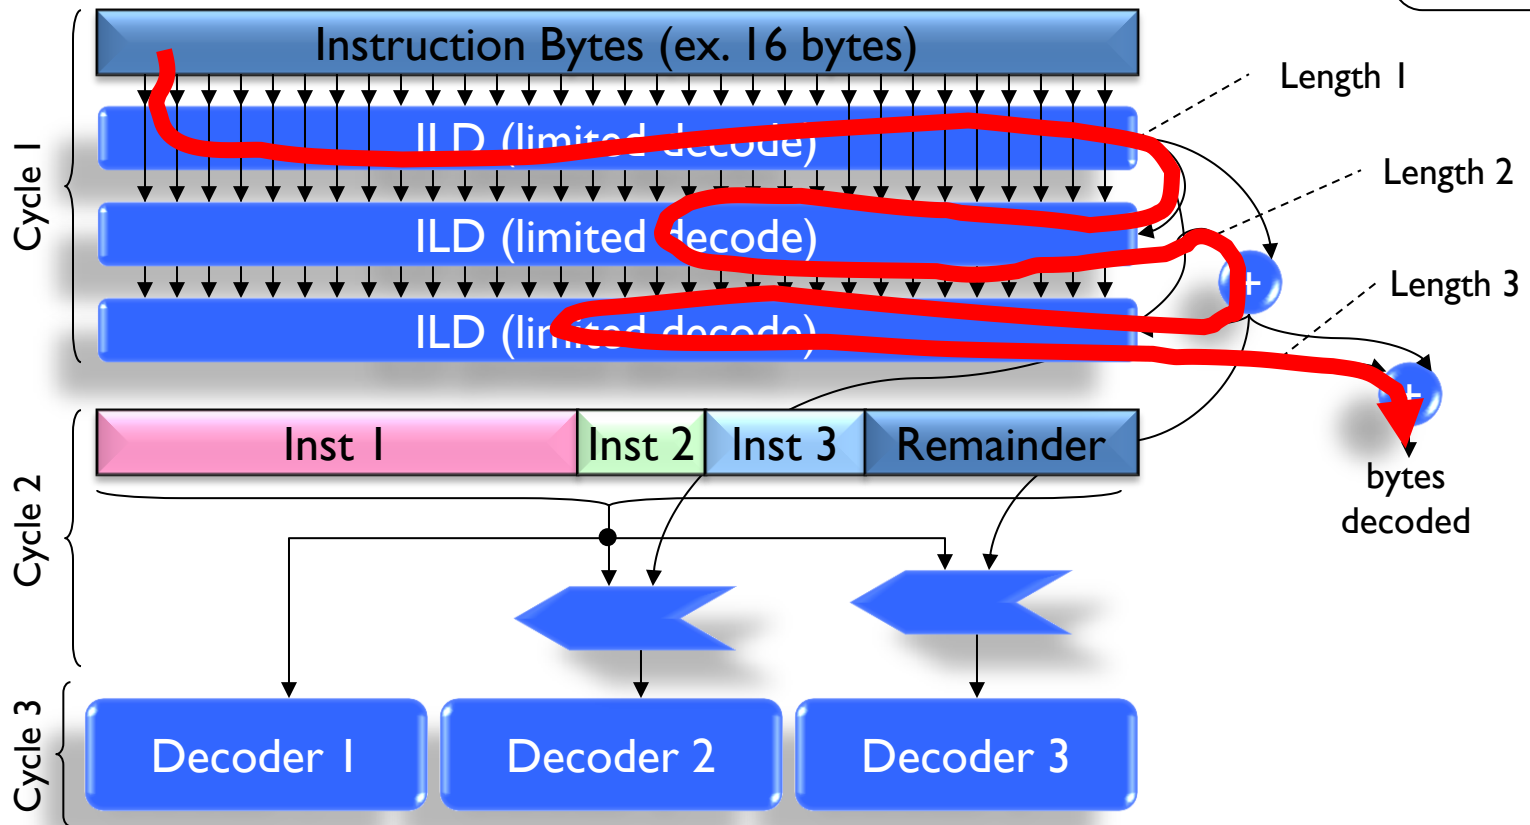
$$PC_{i+1} = PC_i + \text{sizeof}(\text{Mem}[PC_i])$$

$$\begin{aligned} PC_{i+2} &= PC_{i+1} + \text{sizeof}(\text{Mem}[PC_{i+1}]) \\ &= PC_i + \text{sizeof}(\text{Mem}[PC_i]) + \text{sizeof}(\text{Mem}[PC_{i+1}]) \end{aligned}$$

- Can't find start of next insn. before decoding the first
- Must do ILD serially
 - ILD of 4 insns/cycle implies cycle time will be 4x

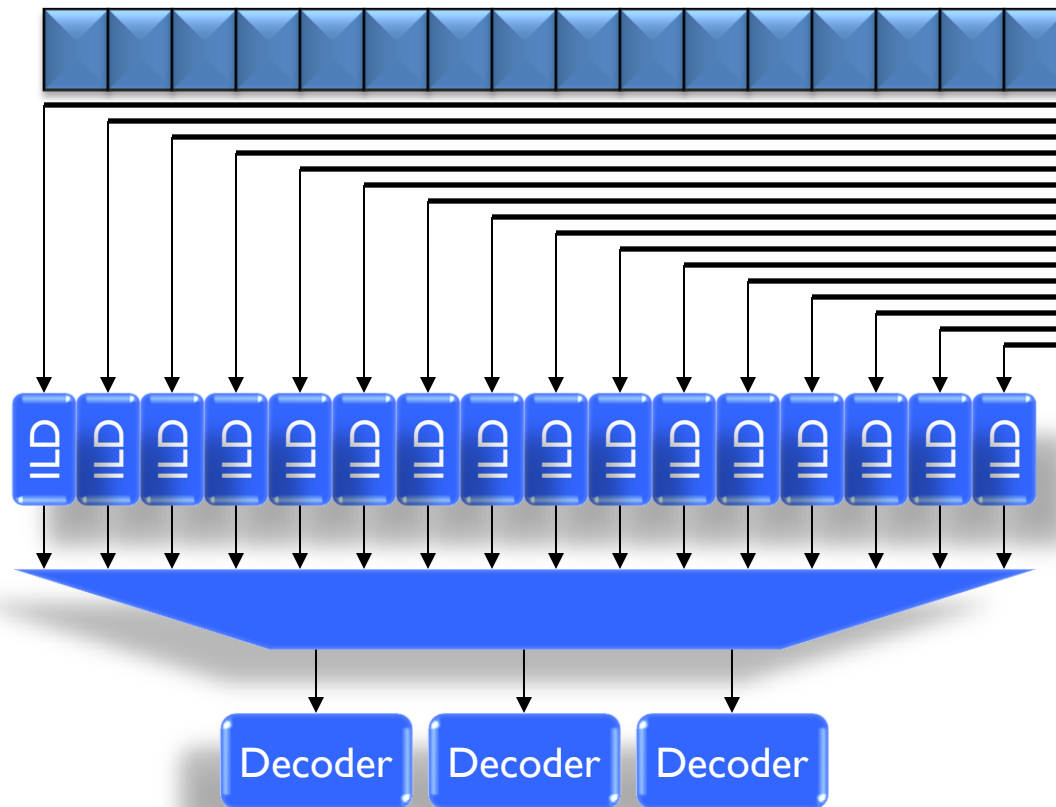
Critical component not pipeline-able

Bad x86 Decode Implementation



ILD dominates cycle time; not scalable

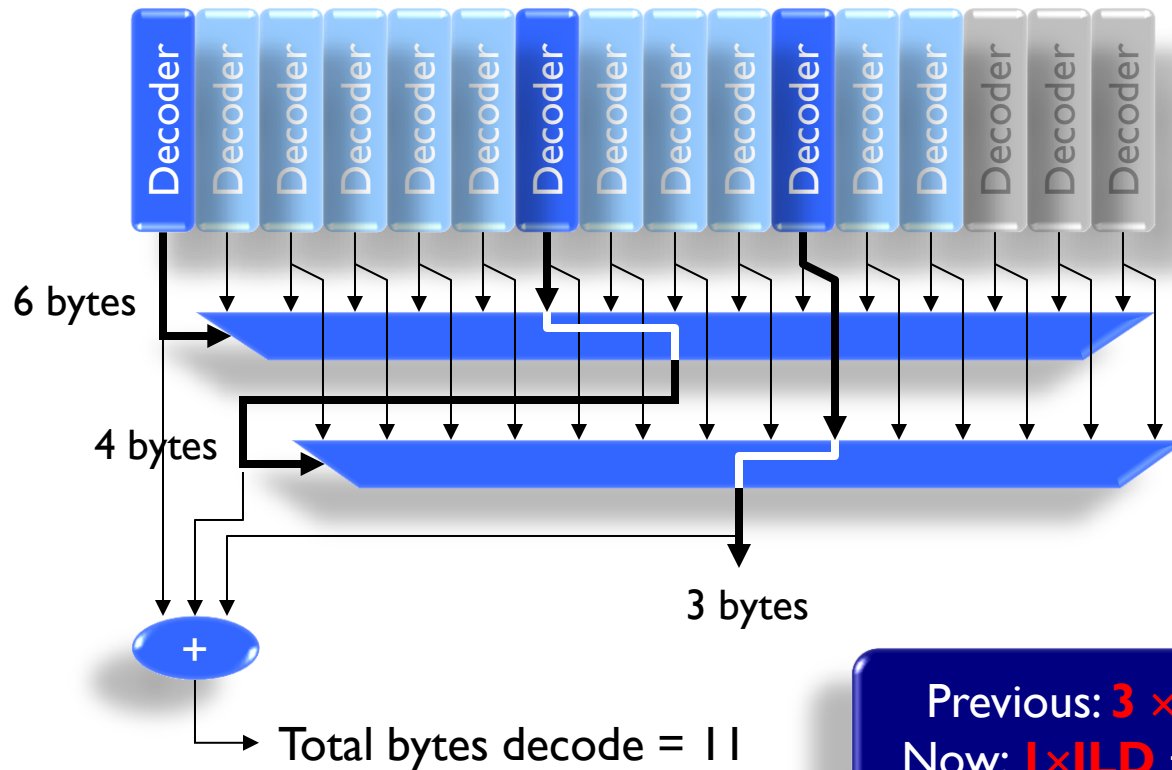
Hardware-Intensive Decode



Decode from every possible instruction starting point!

Giant MUXes to select instruction bytes

ILD in Hardware-Intensive Approach



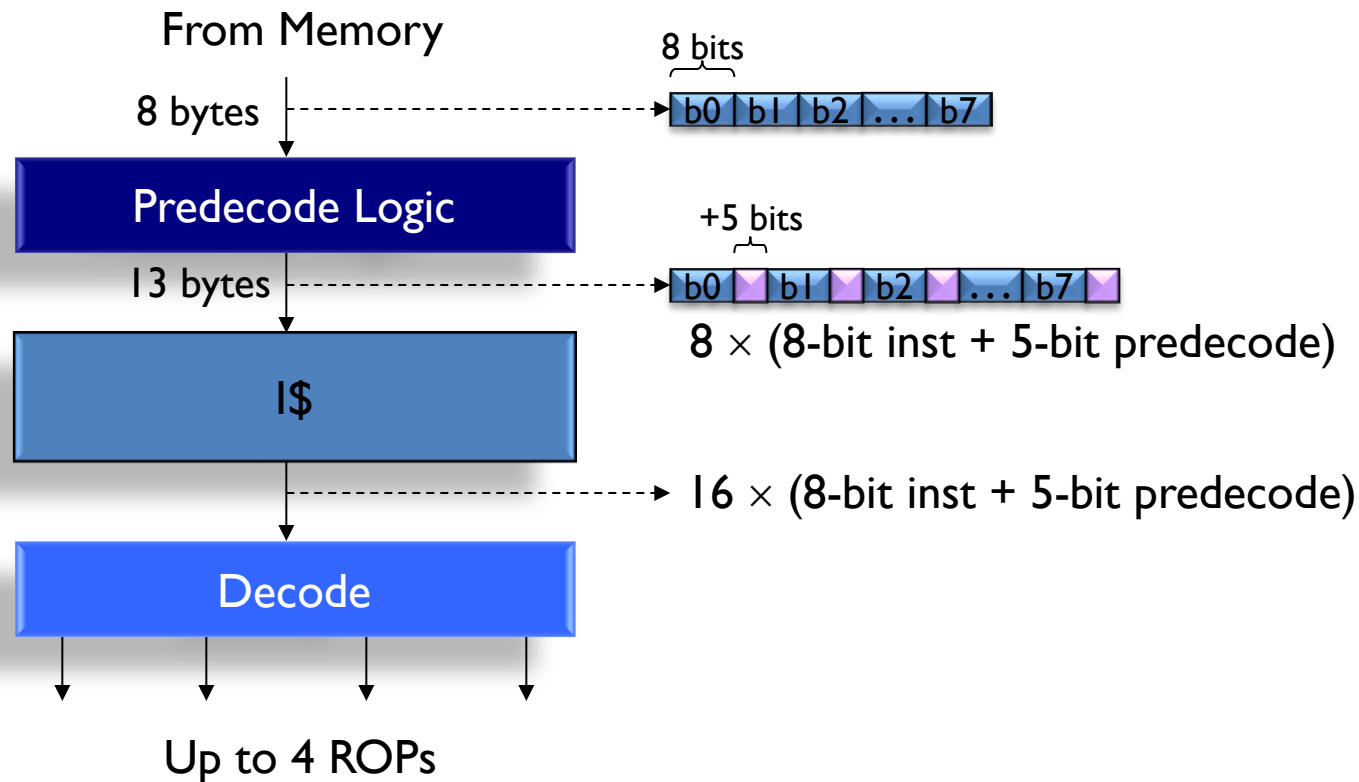
Previous: $3 \times \text{ILD} + 2 \times \text{add}$
 Now: $1 \times \text{ILD} + 2 \times (\text{mux} + \text{add})$

Predecoding

- ILD loop is hardware intensive
 - Impacts latency
 - Consumes substantial power
- If instructions A, B and C are decoded
 - ... lengths for A, B, and C will still be the same next time
 - No need to repeat ILD

Possible to cache the ILD work

Decoder Example: AMD K5 (1/2)

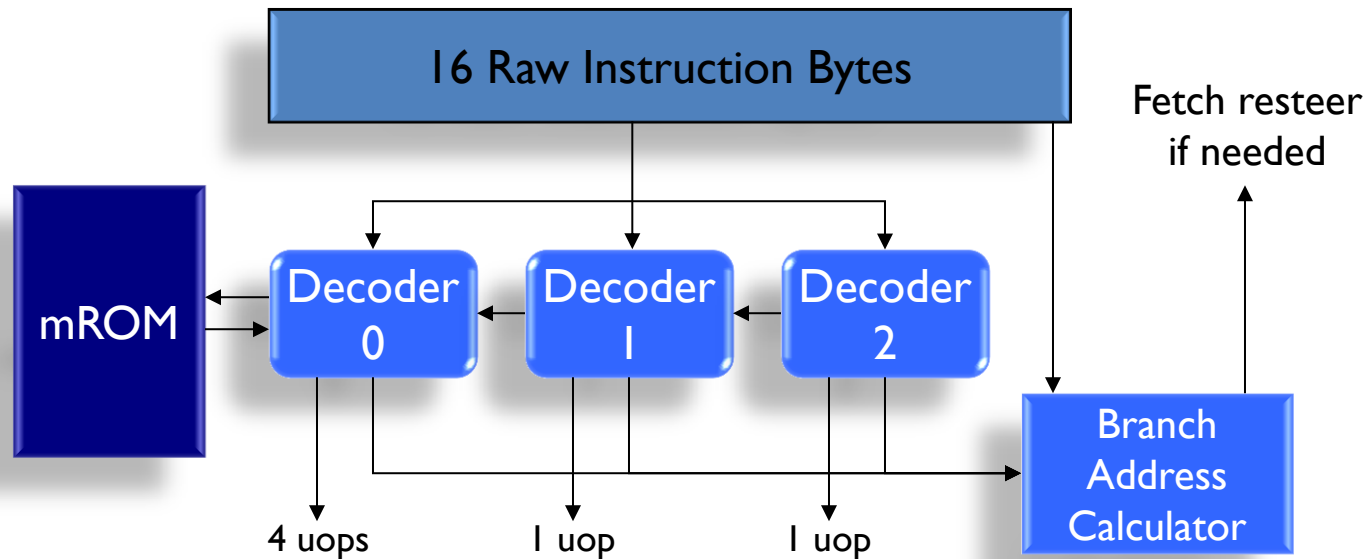


Compute ILD on fetch, store ILD in the I\$

Decoder Example: AMD K5 (2/2)

- Predecode makes decode easier by providing:
 - Instruction start/end location (ILD)
 - Number of ROPs needed per inst
 - Opcode and prefix locations
- Power/performance tradeoffs
 - Larger I\$ (increase data by 62.5%)
 - Longer I\$ latency, more I\$ power consumption
 - Remove logic from decode
 - Shorter pipeline, simpler logic
 - Cache and reused decode work → less decode power
 - Longer effective I-L1 miss latency (ILD on fill)

Decoder Example: Intel P-Pro



- Only Decoder 0 interfaces with the uROM and MS
- If insn. in Decoder 1 or Decoder 2 requires > 1 uop
 - 1) do not generate output
 - 2) shift to Decoder 0 on the next cycle