


# CSE 502: Computer Architecture

Multi-{Socket,Core,Thread}


# Getting More Performance

- Keep pushing IPC and/or frequency
  - Design complexity (time to market)
  - Cooling (cost)
  - Power delivery (cost)
  - ...
- Possible, but too costly

# Bridging the Gap



# Higher Complexity not Worth Effort



# User Visible/Invisible


- All performance gains up to this point were “free”
  - No user intervention required (beyond buying new chip)
    - Recompilation/rewriting could provide even more benefit
  - Higher frequency & higher IPC
  - Same ISA, different micro-architecture
- Multi-processing pushes parallelism above ISA
  - Coarse grained parallelism
    - Provide multiple processing elements
  - User (or developer) responsible for finding parallelism
    - User decides how to use resources

# Sources of (Coarse) Parallelism

- Different applications
  - MP3 player in background while you work in Office
  - Other background tasks: OS/kernel, virus check, etc...
  - Piped applications
    - `gunzip -c foo.gz | grep bar | perl some-script.pl`
- Threads within the same application
  - Java (scheduling, GC, etc...)
  - Explicitly coded multi-threading
    - pthreads, MPI, etc...

# SMP Machines

- SMP = Symmetric Multi-Processing
  - Symmetric = All CPUs have “equal” access to memory
- OS seems multiple CPUs
  - Runs one process (or thread) on each CPU




CPU<sub>0</sub>

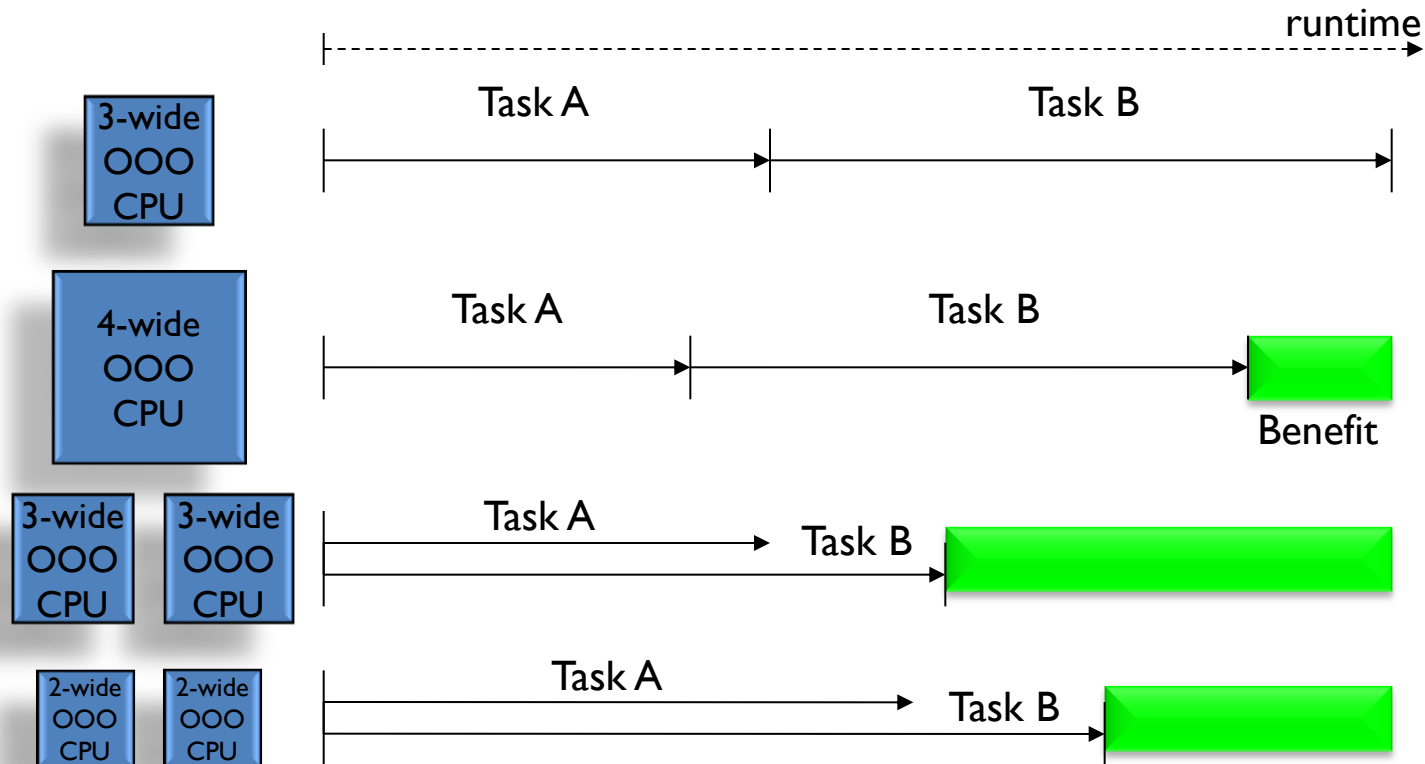
CPU<sub>1</sub>

CPU<sub>2</sub>

CPU<sub>3</sub>



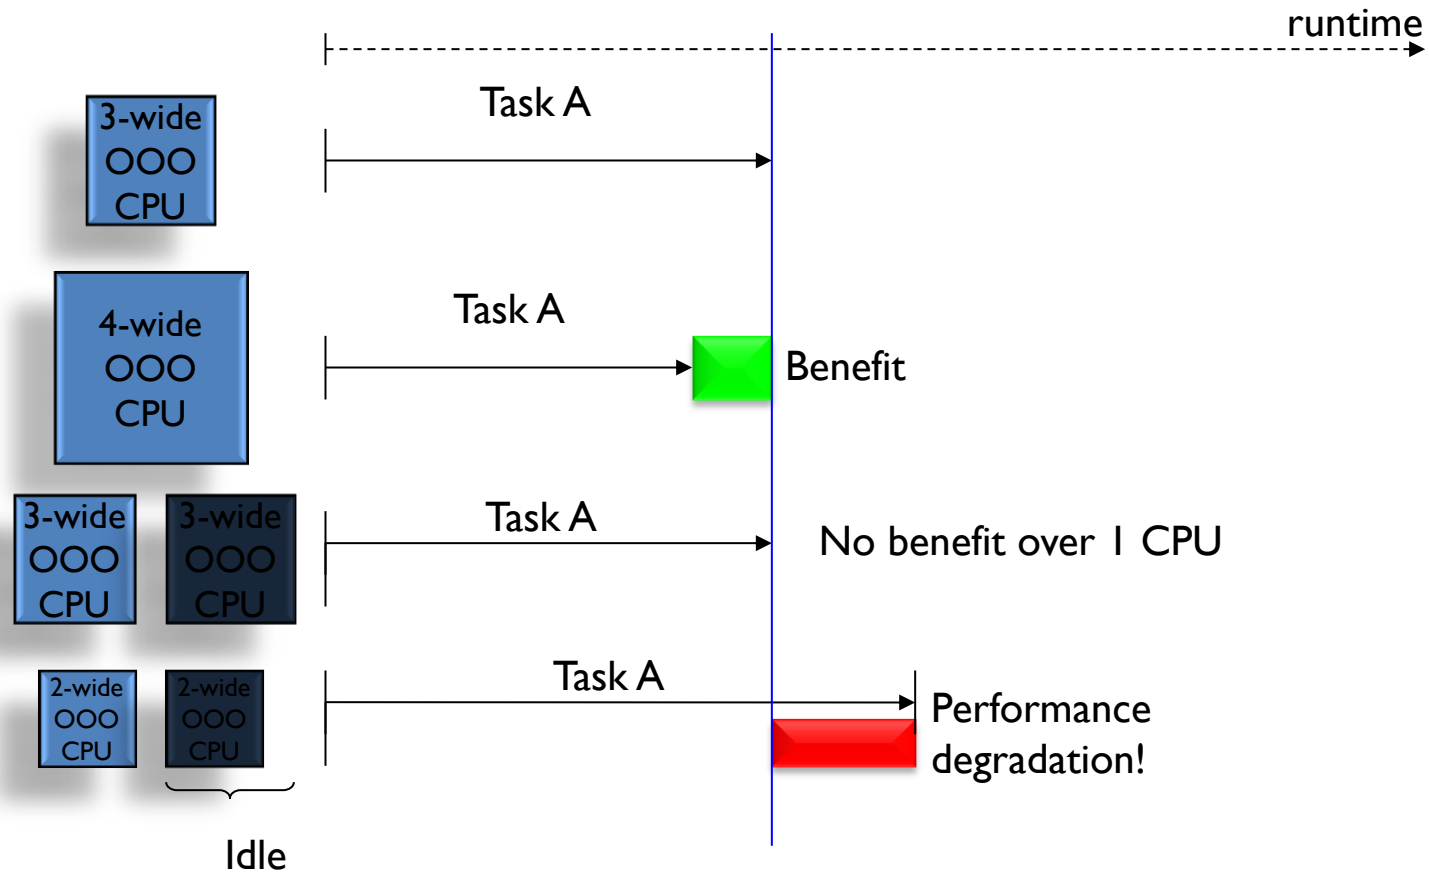
# MP Workload Benefits



Assumes you have multiple tasks/programs to run




# ... If Only One Task Available



# Benefit of MP Depends on Workload

- Limited number of parallel tasks to run on PC
  - Adding more CPUs than tasks provide zero benefit
- For parallel code, Amdahl's law curbs speedup

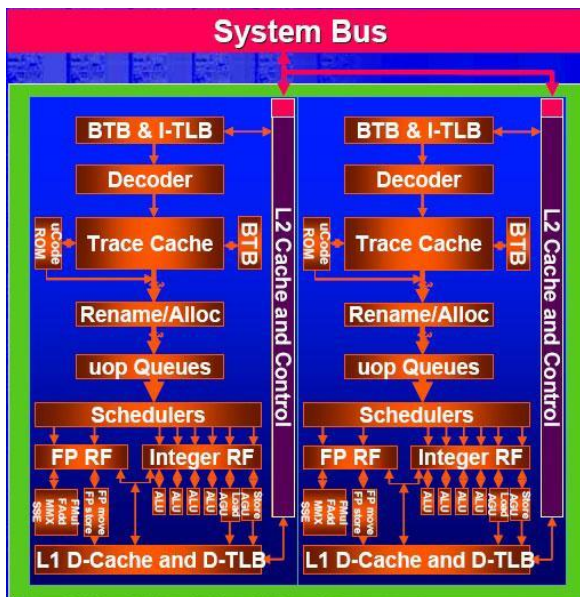


# Hardware Modifications for SMP


- Processor
  - Memory interface
- Motherboard
  - Multiple sockets (one per CPU)
  - Datapaths between CPUs and memory
- Other
  - Case: larger (bigger motherboard, better airflow)
  - Power: bigger power supply for  $N$  CPUs
  - Cooling: more fans to remove  $N$  CPUs' worth of heat

# Chip-Multiprocessing (CMP)

- Simple SMP on the same chip
  - CPUs now called “cores” by hardware designers
  - OS designers still call these “CPUs”



Intel “Smithfield” Block Diagram



AMD Dual-Core Athlon FX

# Benefits of CMP

- Cheaper than multi-chip SMP
  - All/most interface logic integrated on chip
    - Fewer chips
    - Single CPU socket
    - Single interface to memory
  - Less power than multi-chip SMP
    - Communication on-die is more power-efficient than chip-to-chip
- Efficiency
  - Use for transistors instead of wider/more aggressive OoO
  - Potentially better use of hardware resources


# CMP Performance vs. Power

- 2x CPUs not necessarily equal to 2x performance
- 2x CPUs  $\rightarrow$   $\frac{1}{2}$  power for each
  - Maybe a little better than  $\frac{1}{2}$  if resources can be shared
- Back-of-the-Envelope calculation:
  - 3.8 GHz CPU at 100W
  - Dual-core: 50W per CPU
  - $P \propto V^3$ :  $V_{\text{orig}}^3 / V_{\text{CMP}}^3 = 100\text{W} / 50\text{W} \rightarrow V_{\text{CMP}} = 0.8 V_{\text{orig}}$
  - $f \propto V$ :  $f_{\text{CMP}} = 3.0\text{GHz}$

# Multi-Threading

- Uni-Processor: 4-6 wide, lucky if you get 1-2 IPC
  - Poor utilization of transistors
- SMP: 2-4 CPUs, but need independent threads
  - Poor utilization as well (if limited tasks)
- *{Coarse-Grained, Fine-Grained, Simultaneous}-MT*
  - Use single large uni-processor as a multi-processor
    - Core provide multiple hardware contexts (threads)
      - Per-thread PC
      - Per-thread ARF (or map table)
  - Each core appears as multiple CPUs
    - OS designers still call these “CPUs”


# Scalar Pipeline



Dependencies limit functional unit utilization




# Superscalar Pipeline




Higher performance than scalar, but lower utilization

# Chip Multiprocessing (CMP)



Limited utilization when running one thread

# Coarse-Grained Multithreading (1/3)




Only good for long latency ops (i.e., cache misses)

# Coarse-Grained Multithreading (2/3)


- + Sacrifices a little single thread performance
- Tolerates only long latencies (e.g., L2 misses)
- Thread scheduling policy
  - Designate a “preferred” thread (e.g., thread A)
  - Switch to thread B on thread A L2 miss
  - Switch back to A when A L2 miss returns
- Pipeline partitioning
  - None, flush on switch
  - Can’t tolerate latencies shorter than twice pipeline depth
  - Need short in-order pipeline for good performance

# Coarse-Grained Multithreading (3/3)

original pipeline




thread scheduler



L2 miss?

# Fine-Grained Multithreading (1/3)



Saturated workload  $\rightarrow$  Lots of threads



Unsaturated workload  $\rightarrow$  Lots of stalls


Intra-thread dependencies still limit performance

# Fine-Grained Multithreading (2/3)

- Sacrifices significant single-thread performance
- + Tolerates everything
  - + L2 misses
  - + Mispredicted branches
  - + etc...
- Thread scheduling policy
  - Switch threads often (e.g., every cycle)
  - Use round-robin policy, skip threads with long-latency ops
- Pipeline partitioning
  - Dynamic, no flushing
  - Length of pipeline doesn't matter


# Fine-Grained Multithreading (3/3)

- (Many) more threads
- Multiple threads in pipeline at once





# Simultaneous Multithreading (1/3)





Max utilization of functional units

# Simultaneous Multithreading (2/3)

- + Tolerates all latencies
- ± Sacrifices some single thread performance
- Thread scheduling policy
  - Round-robin (like Fine-Grained MT)
- Pipeline partitioning
  - Dynamic
- Examples
  - Pentium4 (hyper-threading): 5-way issue, 2 threads
  - Alpha 21464: 8-way issue, 4 threads (canceled)

# Simultaneous Multithreading (3/3)



# Issues for SMT

- Cache interference
  - Concern for all MT variants
  - Shared memory SPMD threads help here
    - Same insns. → share I\$
    - Shared data → less D\$ contention
    - MT is good for “server” workloads
  - SMT might want a larger L2 (which is OK)
    - Out-of-order tolerates L1 misses
- Large map table and physical register file
  - $\text{\#maptable-entries} = (\text{\#threads} * \text{\#arch-regs})$
  - $\text{\#phys-regs} = (\text{\#threads} * \text{\#arch-regs}) + \text{\#in-flight insns}$

# Latency vs Throughput

- MT trades (single-thread) latency for throughput
  - Sharing processor degrades latency of individual threads
  - But improves aggregate latency of both threads
  - Improves utilization
- Example
  - Thread A: individual latency=10s, latency with thread B=15s
  - Thread B: individual latency=20s, latency with thread A=25s
  - Sequential latency (first A then B or vice versa): 30s
  - Parallel latency (A and B simultaneously): 25s
  - MT slows each thread by 5s
  - But improves total latency by 5s

Benefits of MT depend on workload

# CMP vs MT

- If you wanted to run multiple threads would you build a...
  - Chip multiprocessor (CMP): multiple separate pipelines?
  - A multithreaded processor (MT): a single larger pipeline?
- Both will get you throughput on multiple threads
  - CMP will be simpler, possibly faster clock
  - SMT will get you better performance (IPC) on a single thread
    - SMT is basically an ILP engine that converts TLP to ILP
    - CMP is mainly a TLP engine
- Do both (CMP or MTs), e.g., Sun UltraSPARC T1
  - 8 processors, each with 4-threads (fine-grained threading)
  - 1Ghz clock, in-order, short pipeline
  - Designed for power-efficient “throughput computing”

# Combining MP Techniques (1/2)

- System can have SMP, CMP, and SMT at the same time
- Example machine with 32 threads
  - Use 2-socket SMP motherboard with two chips
  - Each chip with an 8-core CMP
  - Where each core is 2-way SMT
- Makes life difficult for the OS scheduler
  - OS needs to know which CPUs are...
    - Real physical processor (SMP): highest independent performance
    - Cores in same chip: fast core-to-core comm., but shared resources
    - Threads in same core: competing for resources
  - Distinct apps. scheduled on different CPUs
  - Cooperative apps. (e.g., pthreads) scheduled on same core
  - Use SMT as last choice (or don't use for some apps.)

# Combining MP Techniques (2/2)

Softech.cz



Softech.cz




# Scalability Beyond the Machine



# Server Racks



# Datacenters (1/2)





# Datacenters (2/2)

